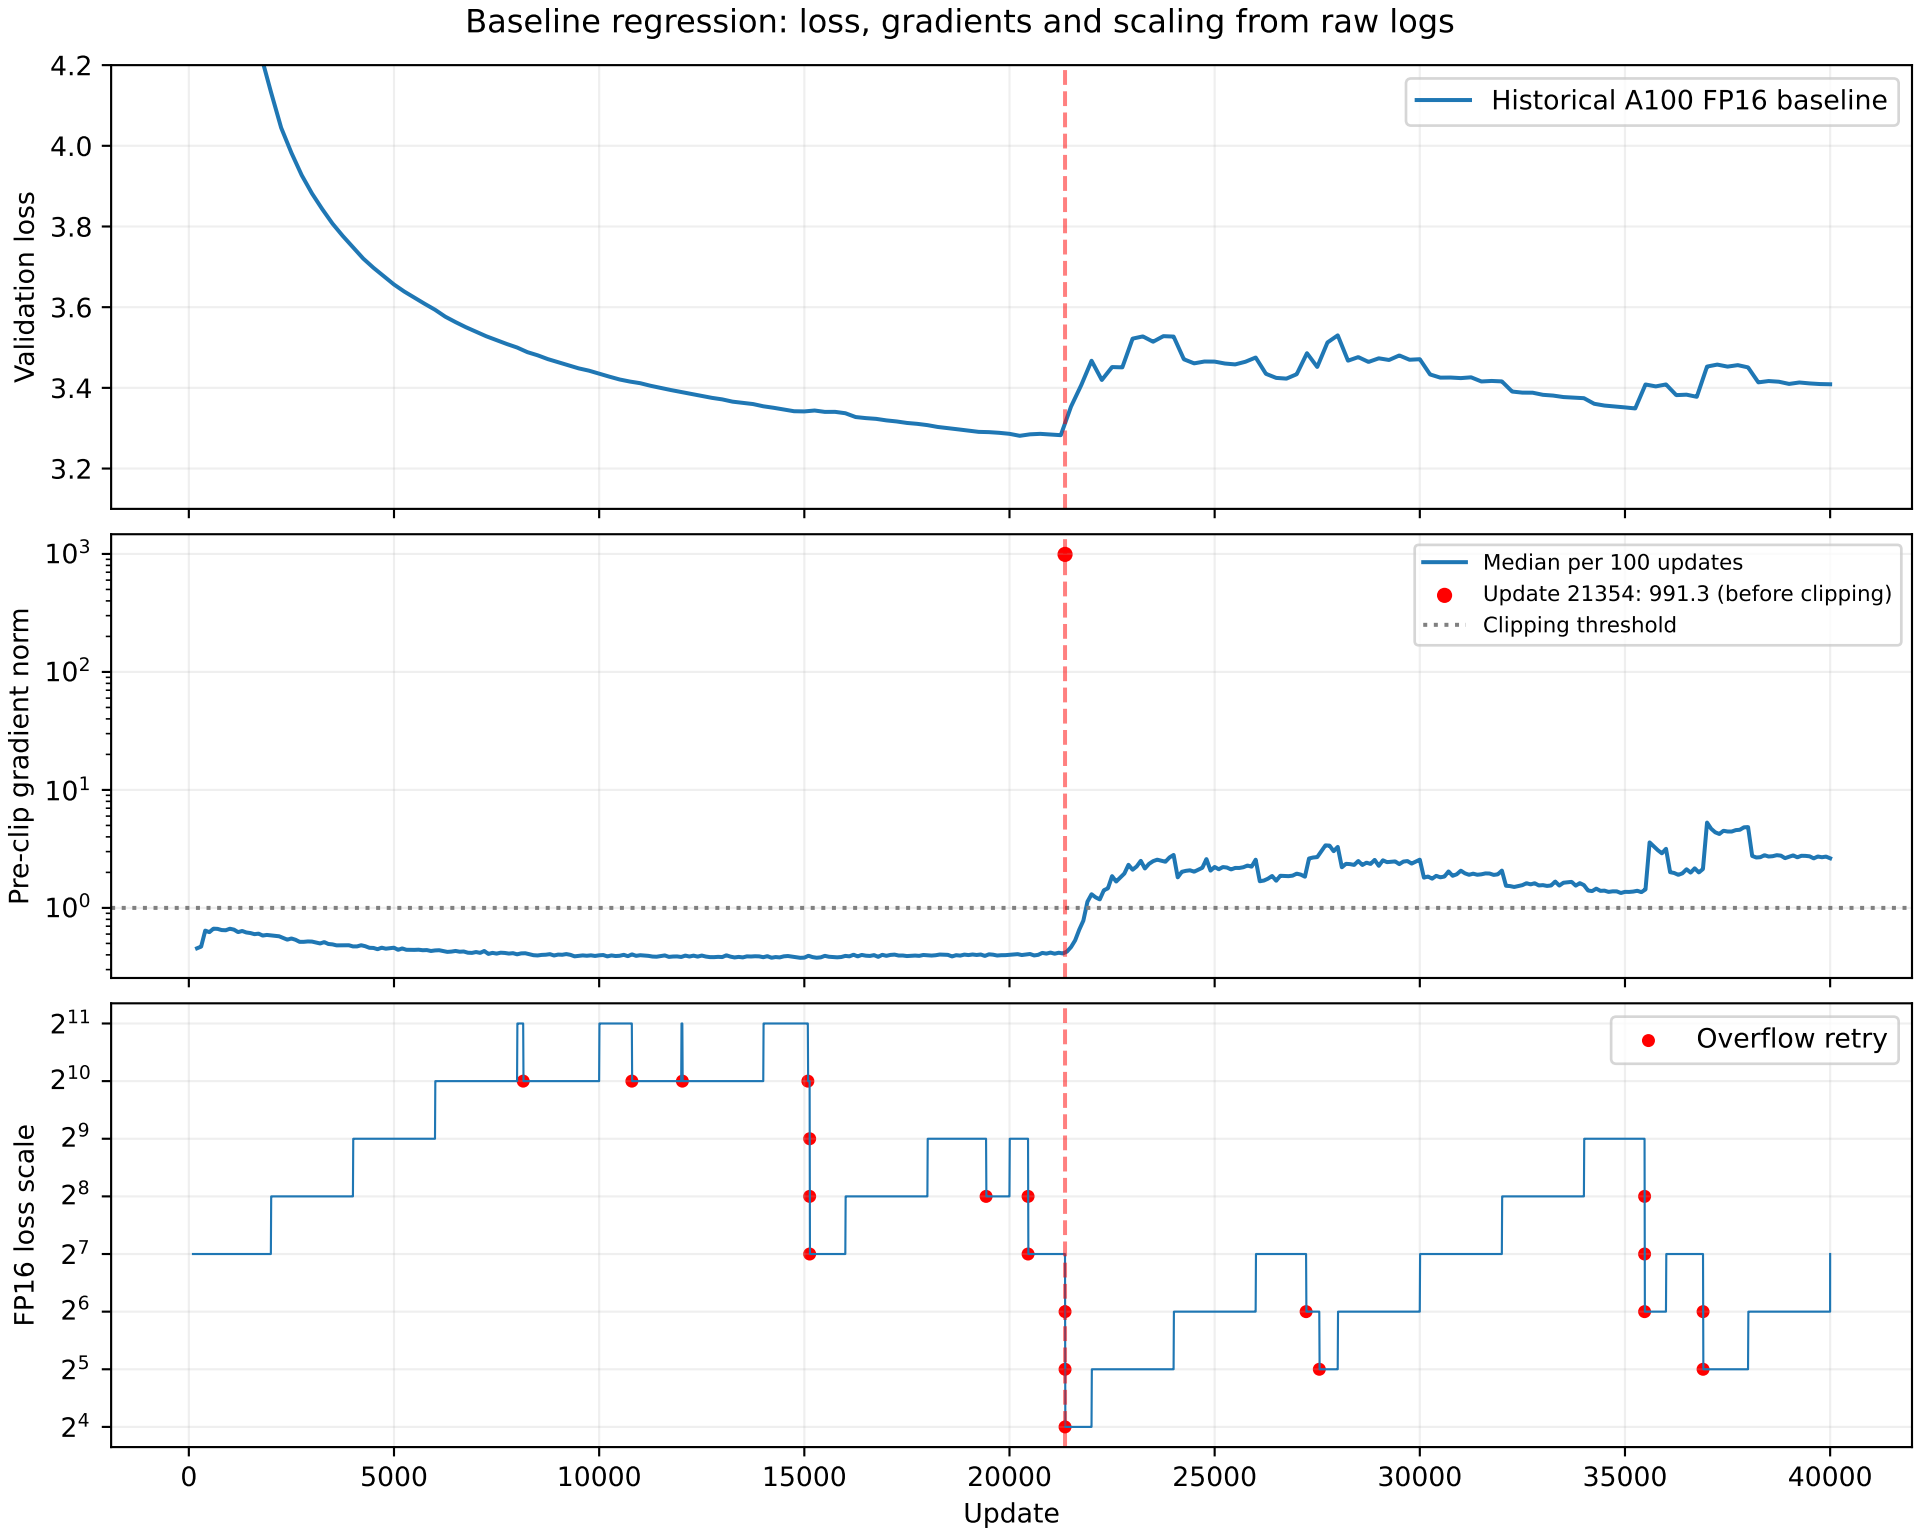
Baseline regression: loss, gradients and scaling from raw logs
Validation loss
4.2
4.0
3.8
3.6
3.4
3.2
Historical A100 FP16 baseline
Pre-clip gradient norm
10^3
10^2
10^1
10^0
Median per 100 updates
Update 21354: 991.3 (before clipping)
Clipping threshold
FP16 loss scale
2^{11}
2^{10}
2^9
2^8
2^7
2^6
2^5
2^4
Overflow retry
0
5000
10000
15000
20000
25000
30000
35000
40000
Update