Eight downstream tasks: scores and differences

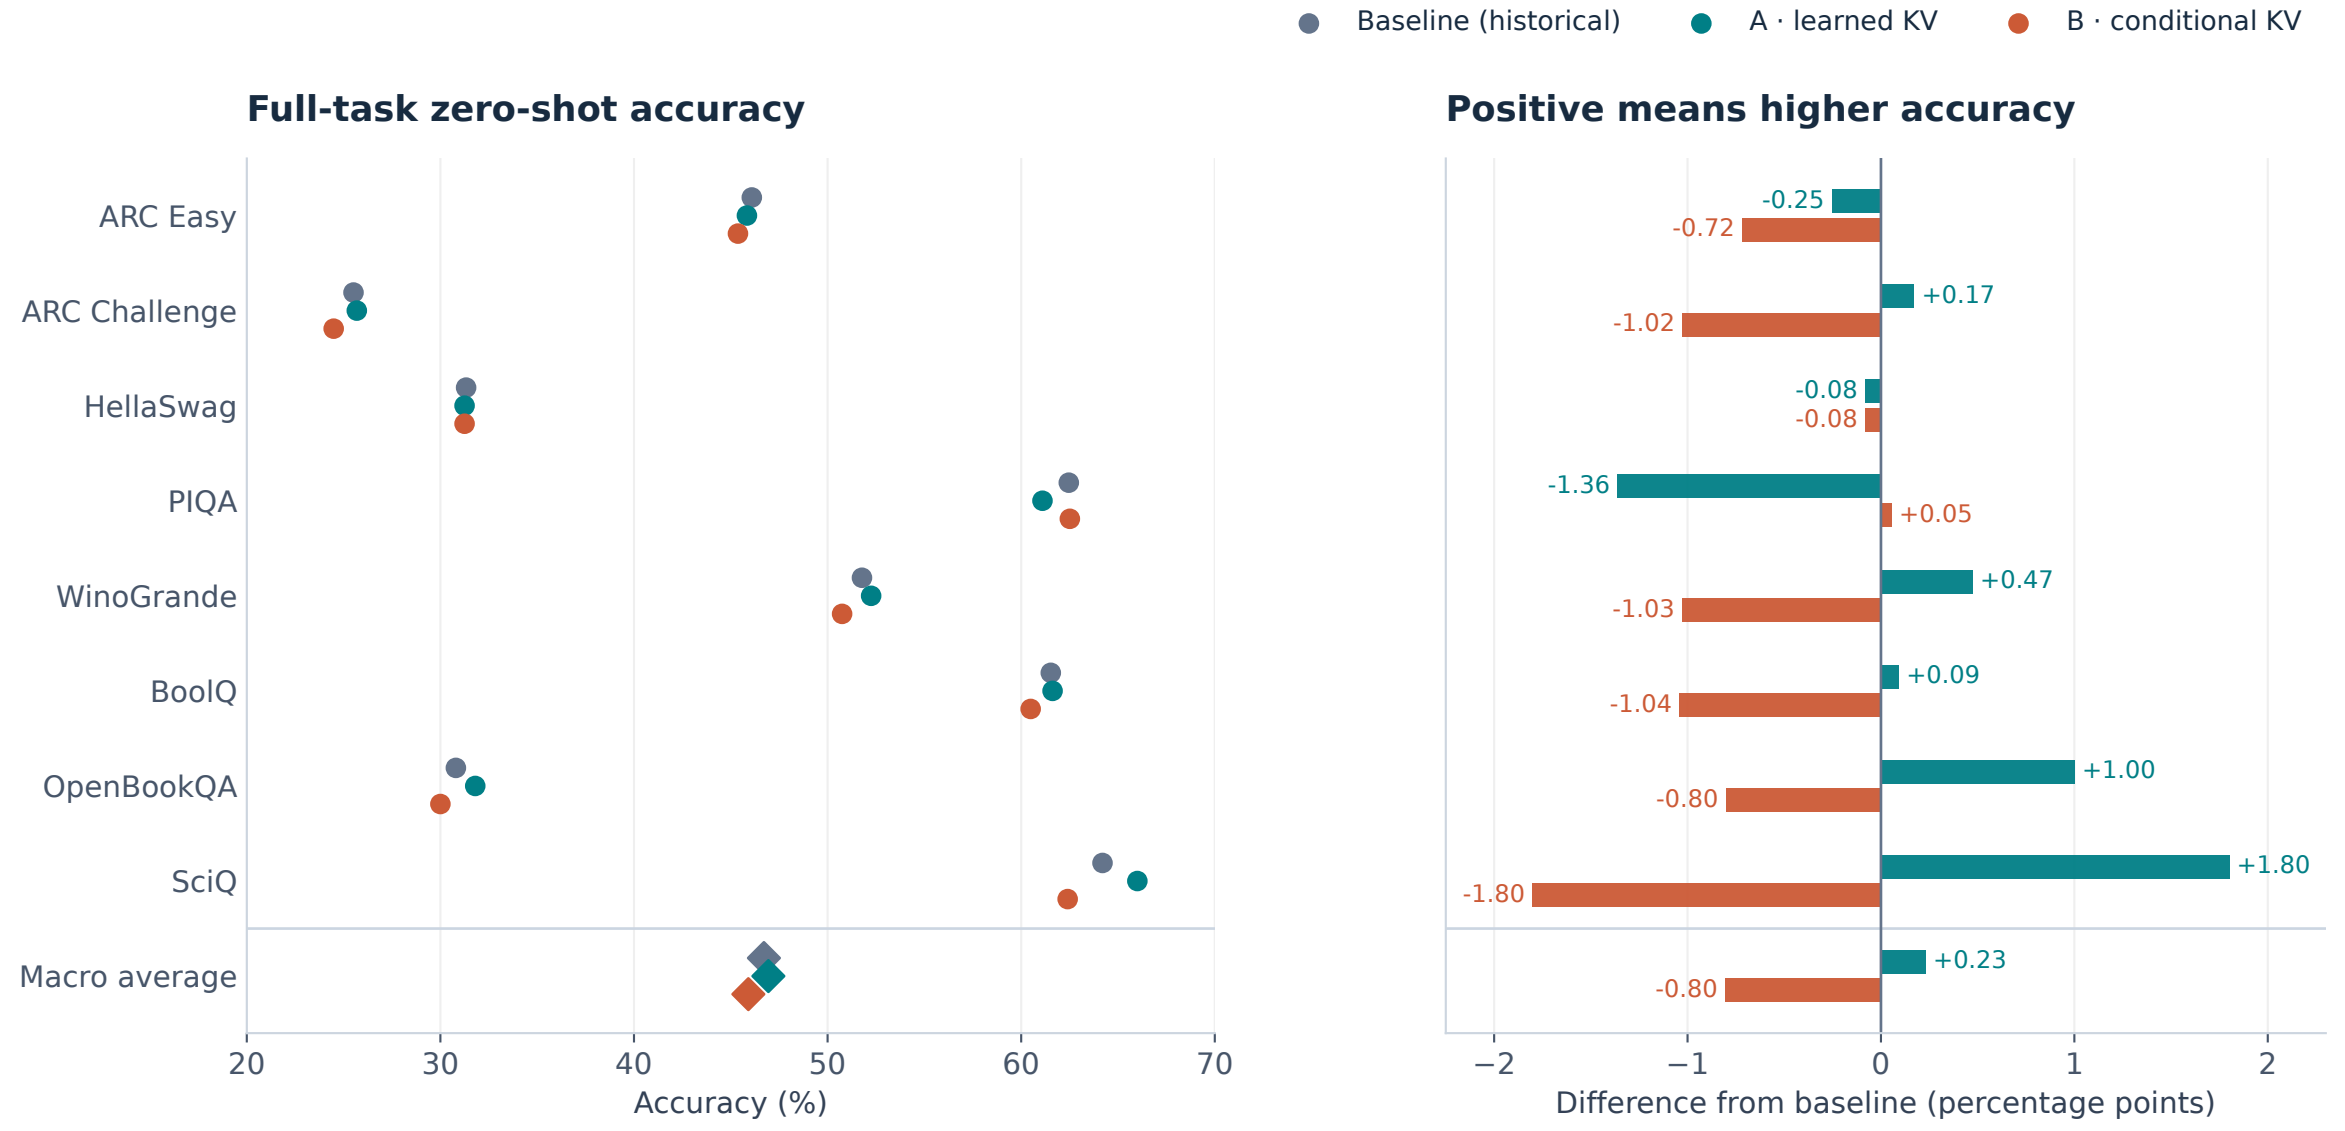

21,465 examples per A/B arm; no limit. BoolQ / WinoGrande: acc; other tasks: acc_norm.
A − baseline: +0.23 pp; B − baseline: −0.80 pp. Baseline differences are descriptive, not paired significance tests.