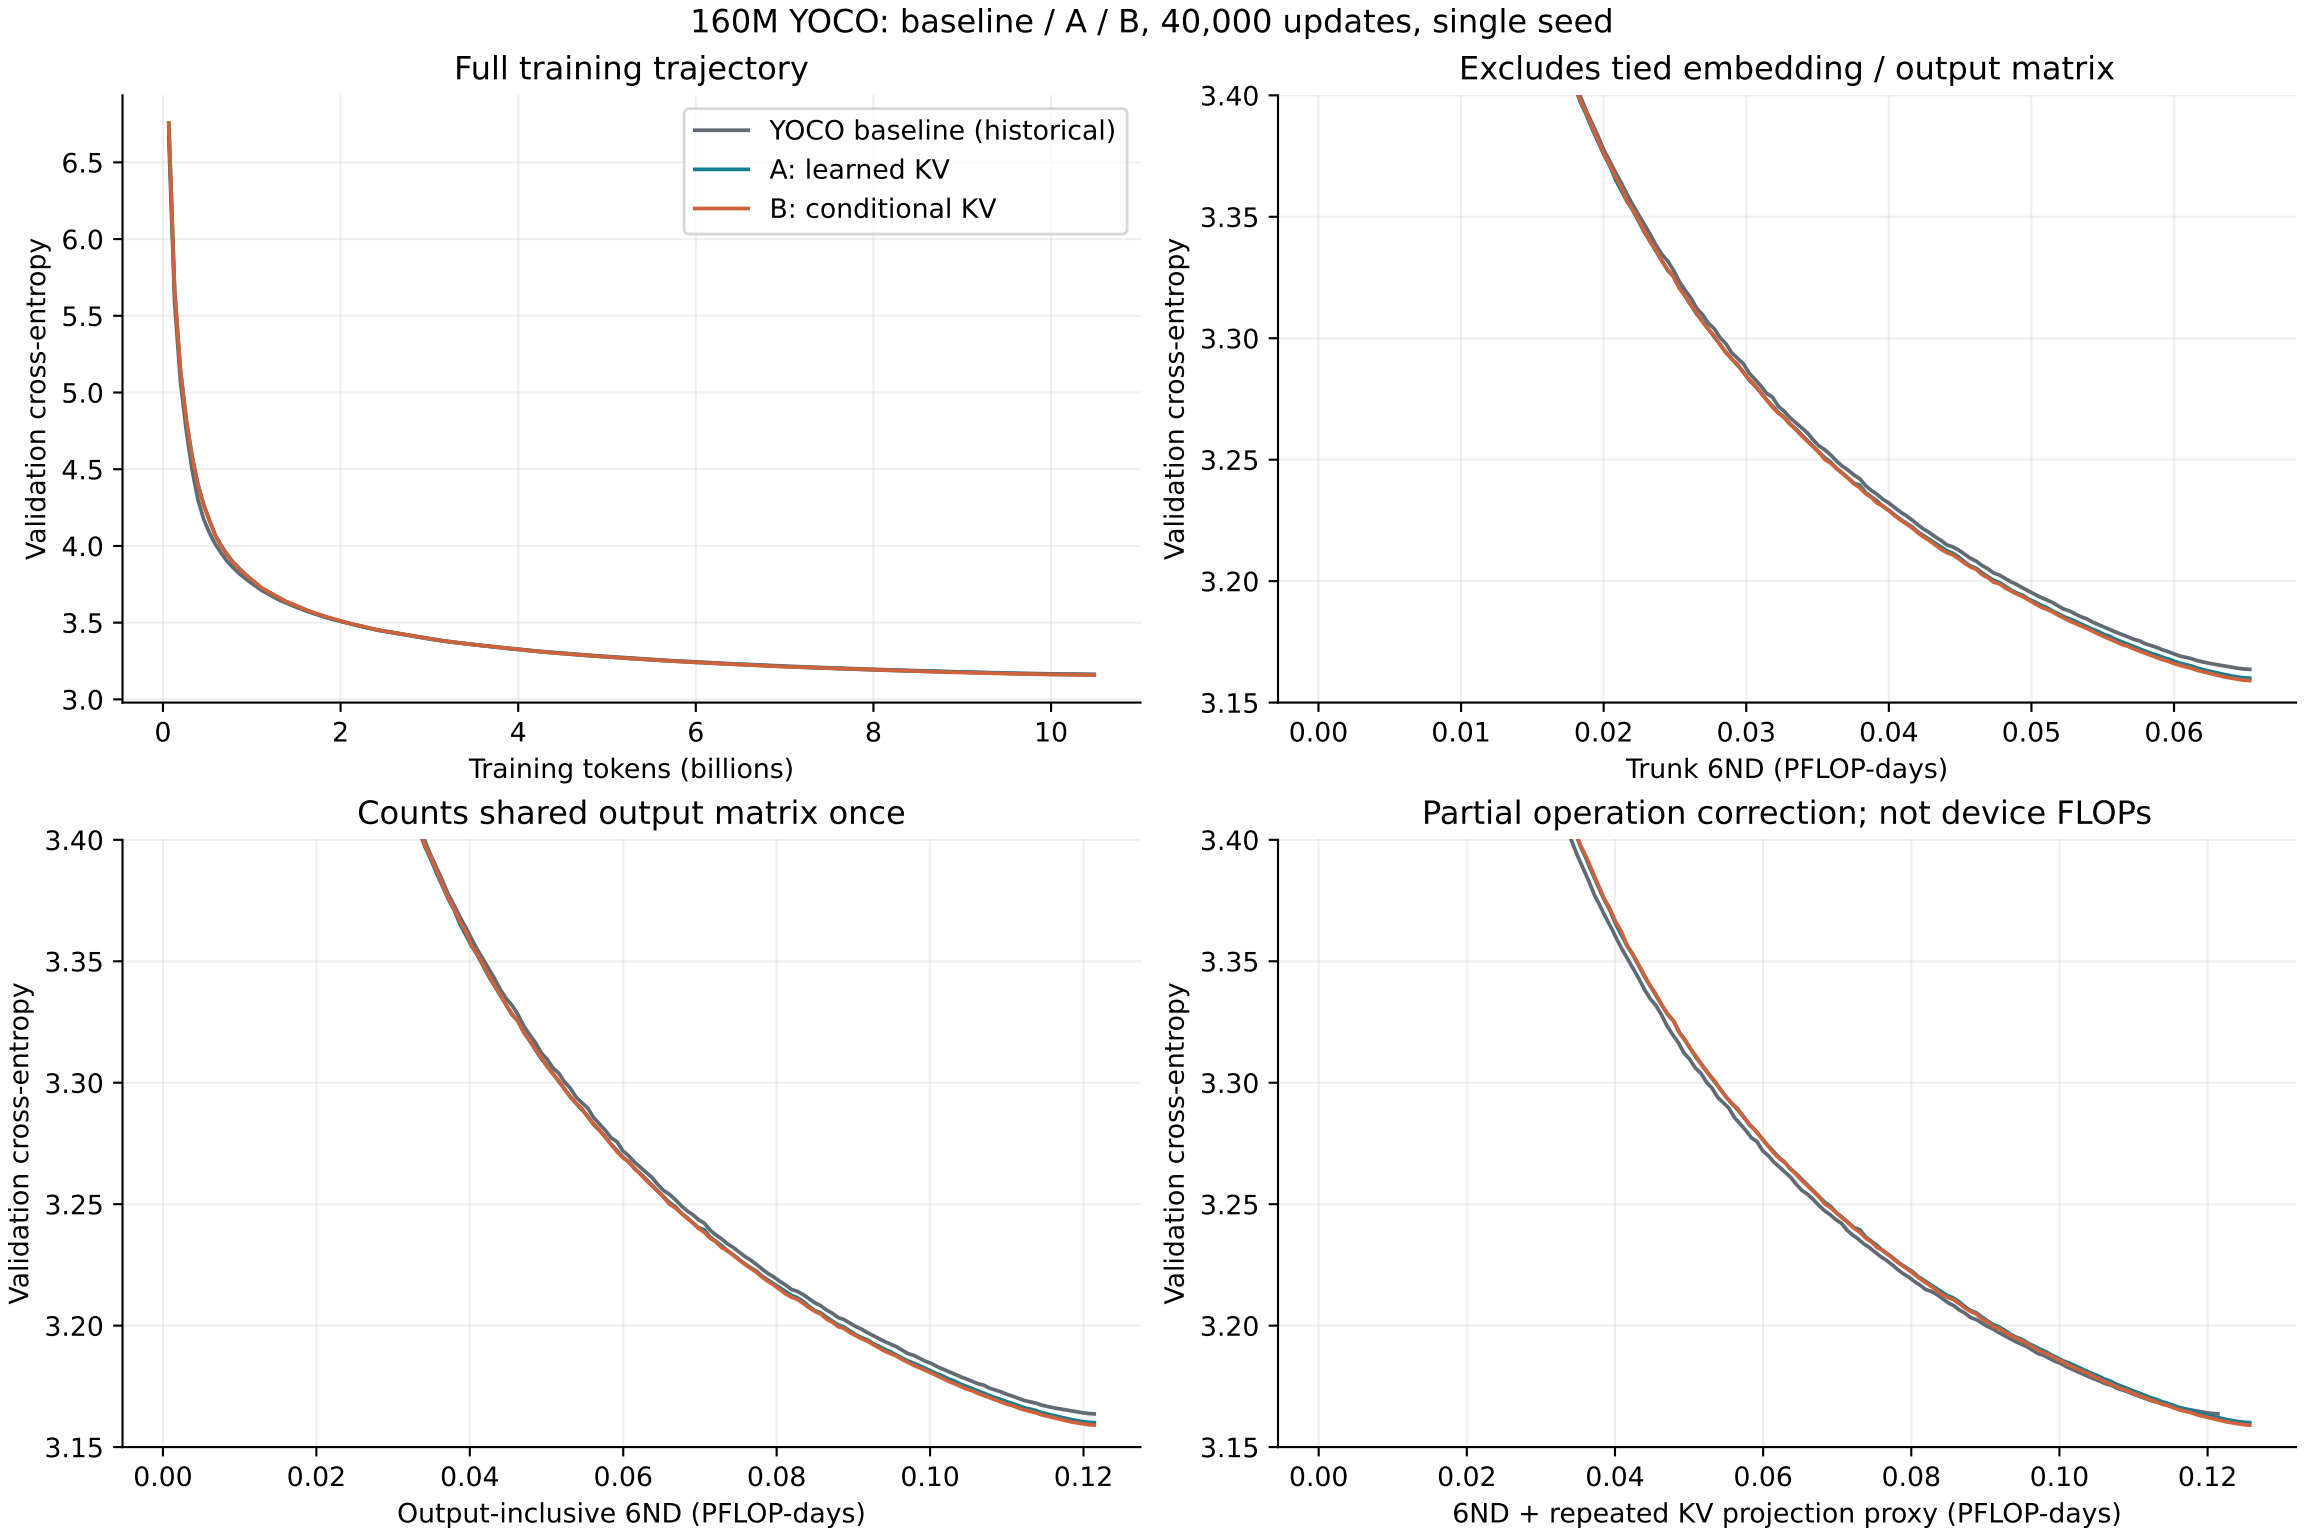
160M YOCO: baseline / A / B, 40,000 updates, single seed
Full training trajectory
YOCO baseline (historical)
A: learned KV
B: conditional KV
Validation cross-entropy
Training tokens (billions)
Excludes tied embedding / output matrix
Trunk 6ND (PFLOP-days)
Counts shared output matrix once
Output-inclusive 6ND (PFLOP-days)
Partial operation correction; not device FLOPs
6ND + repeated KV projection proxy (PFLOP-days)