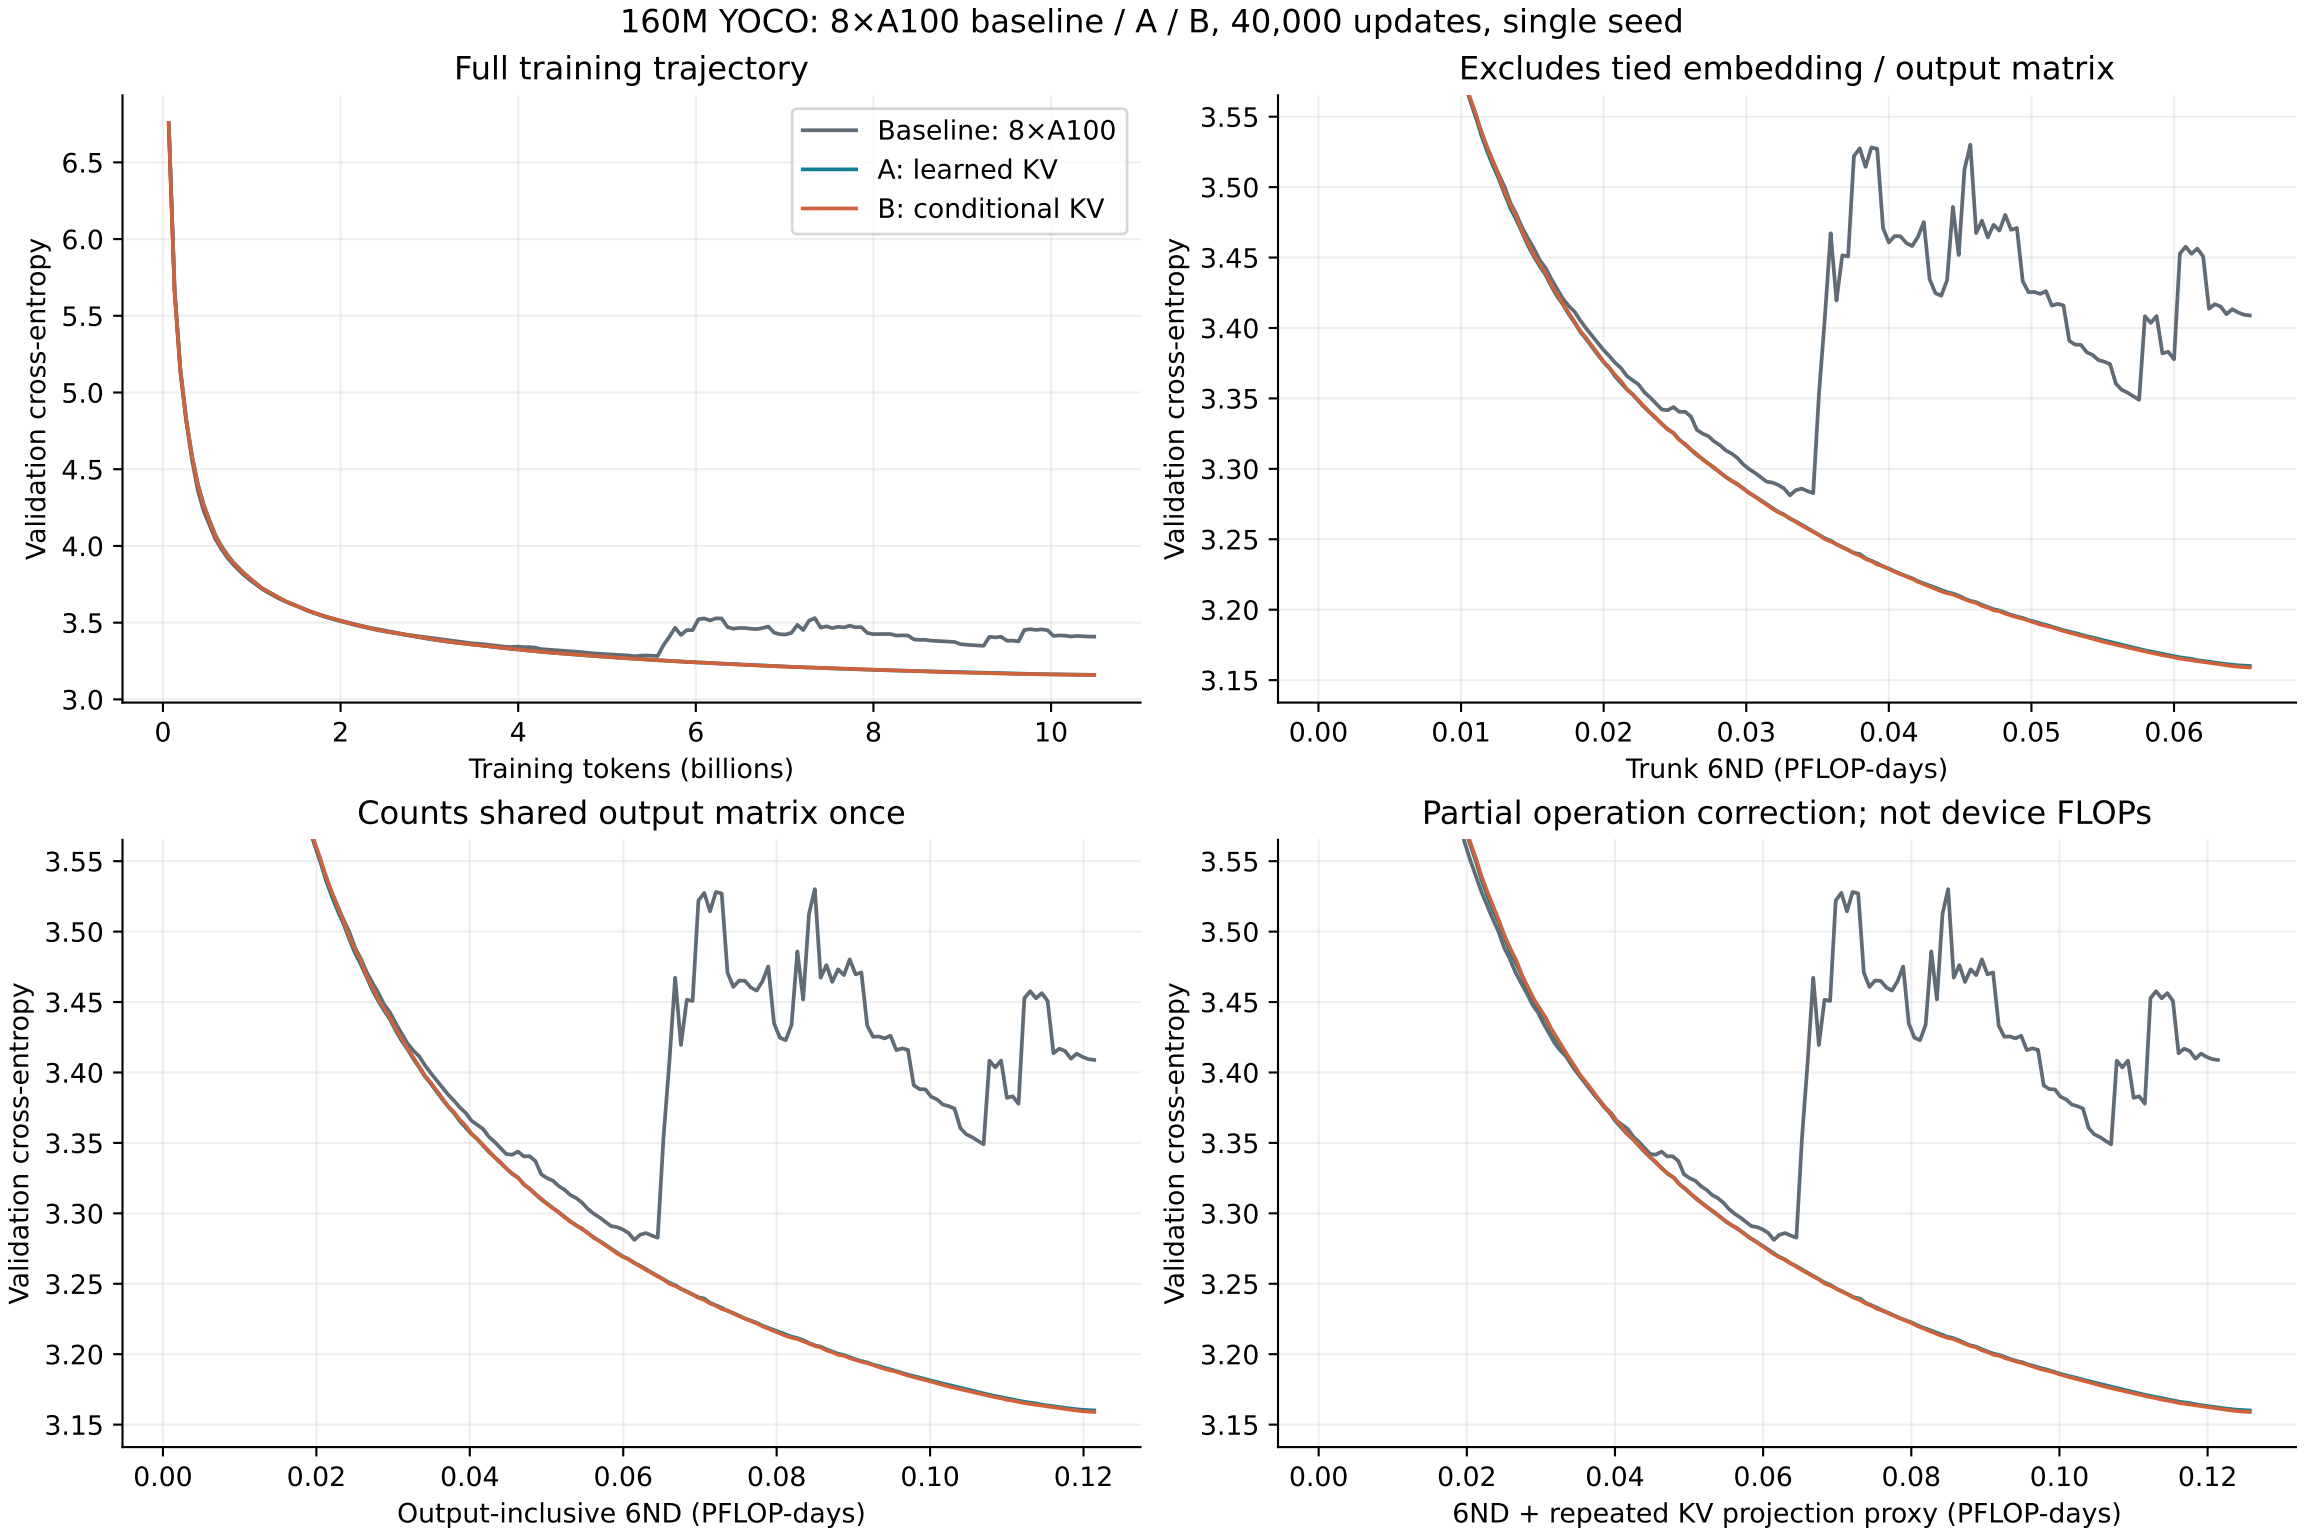
160M YOCO: 8×A100 baseline / A / B, 40,000 updates, single seed
Full training trajectory
Excludes tied embedding / output matrix
Counts shared output matrix once
Partial operation correction; not device FLOPs
Baseline: 8×A100
A: learned KV
B: conditional KV
Validation cross-entropy
Training tokens (billions)
Trunk 6ND (PFLOP-days)
Output-inclusive 6ND (PFLOP-days)
6ND + repeated KV projection proxy (PFLOP-days)