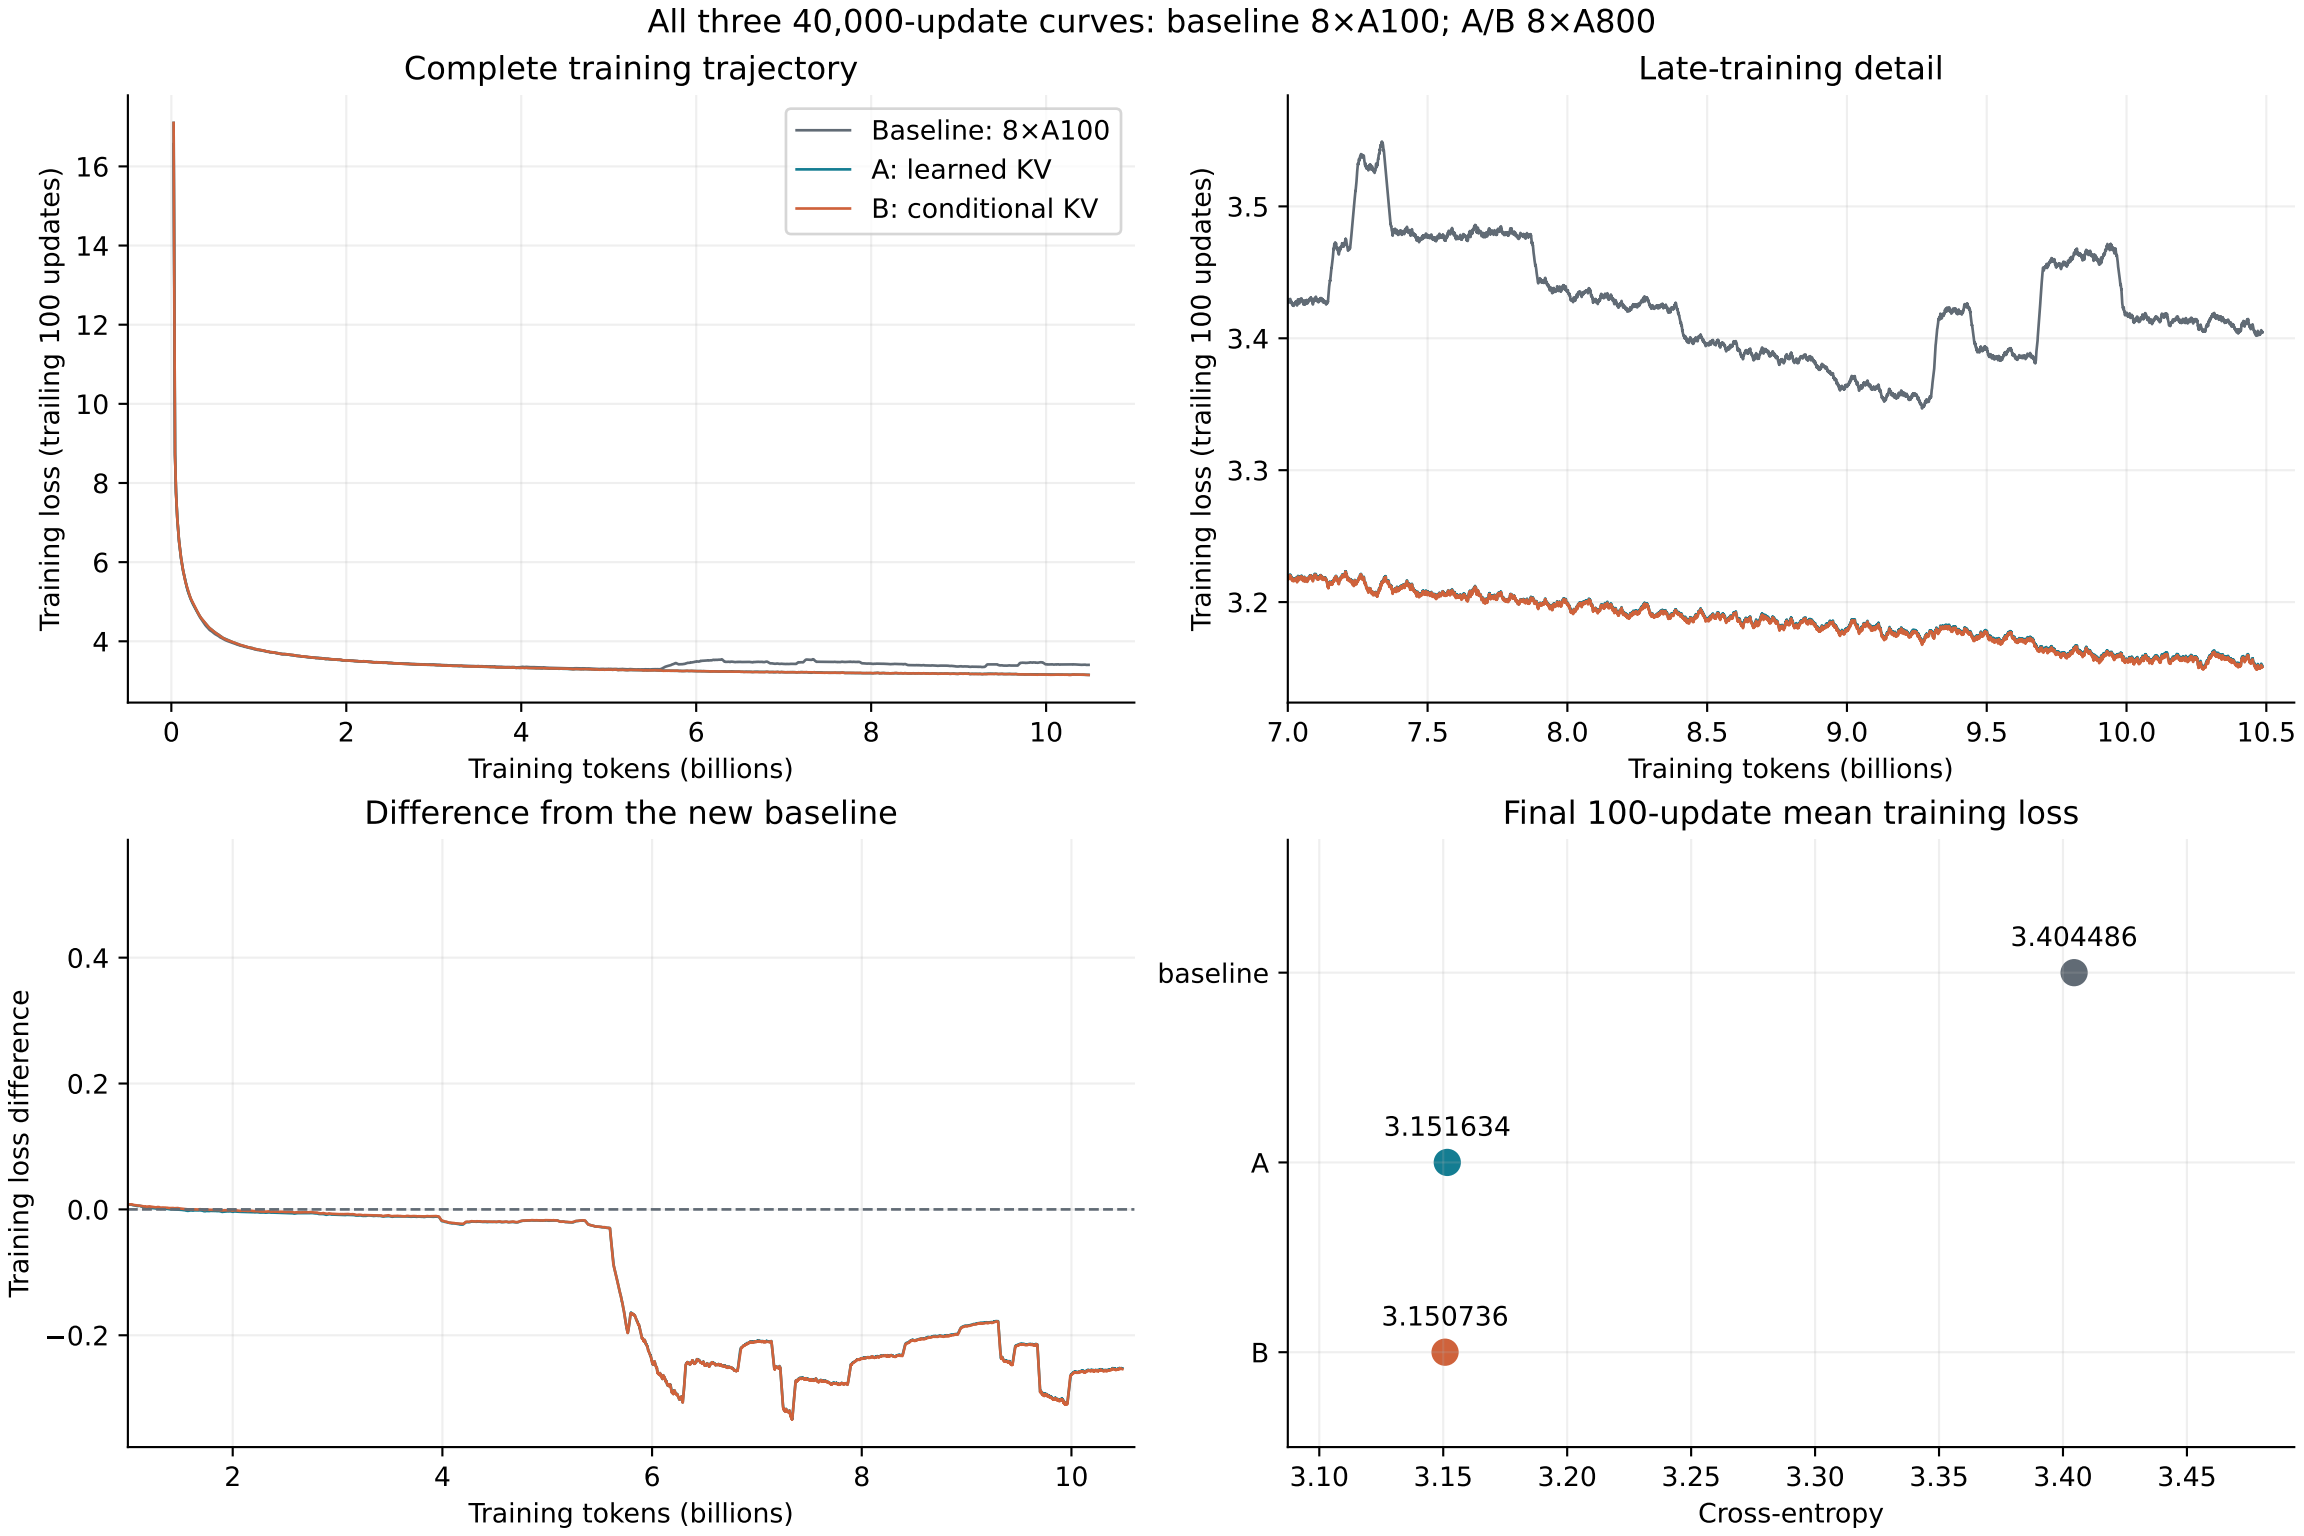

All three 40,000-update curves: baseline 8×A100; A/B 8×A800
Complete training trajectory
Baseline: 8×A100
A: learned KV
B: conditional KV
Training loss (trailing 100 updates)
Training tokens (billions)
Late-training detail
Training loss (trailing 100 updates)
Training tokens (billions)
Difference from the new baseline
Training loss difference
Training tokens (billions)
Final 100-update mean training loss
baseline
A
B
3.404486
3.151634
3.150736
Cross-entropy