YOCO vs Transformer: 2K Quality and V100 Inference Efficiency

Official lm-eval tasks (zero-shot, full test sets)

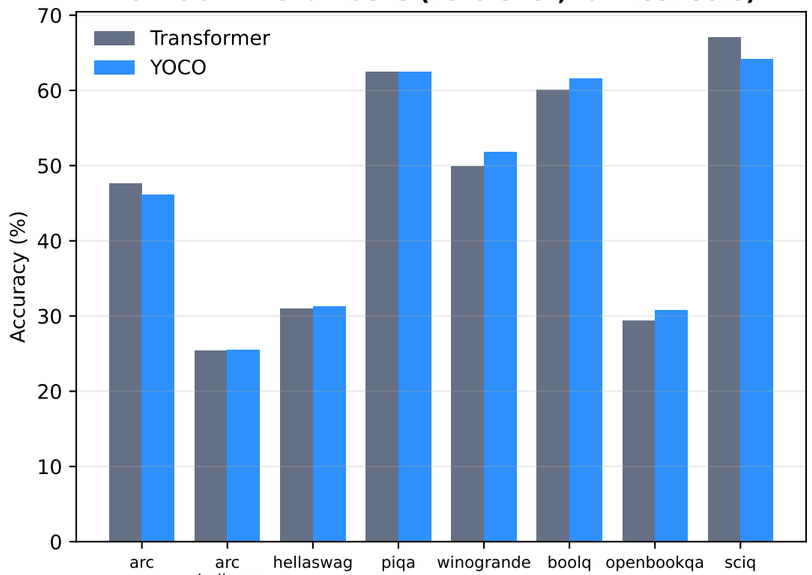

Fixed 2048 prompt + 64-token generation

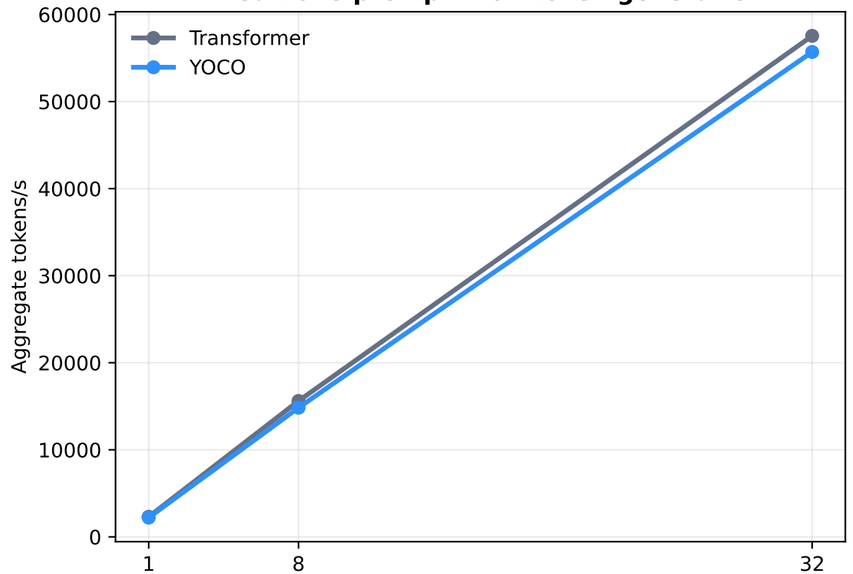

32-GPU fixed-shape speedup

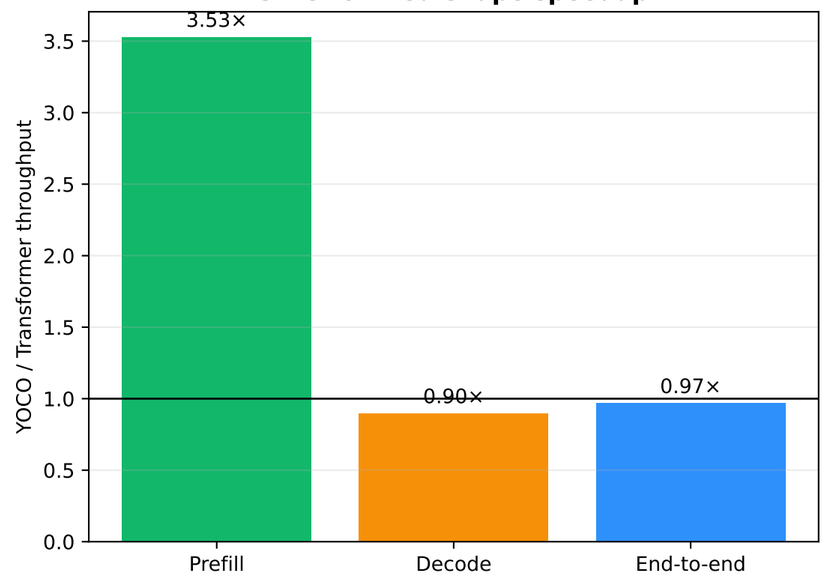

Dynamic-shape trace: 8 prompt/generation patterns

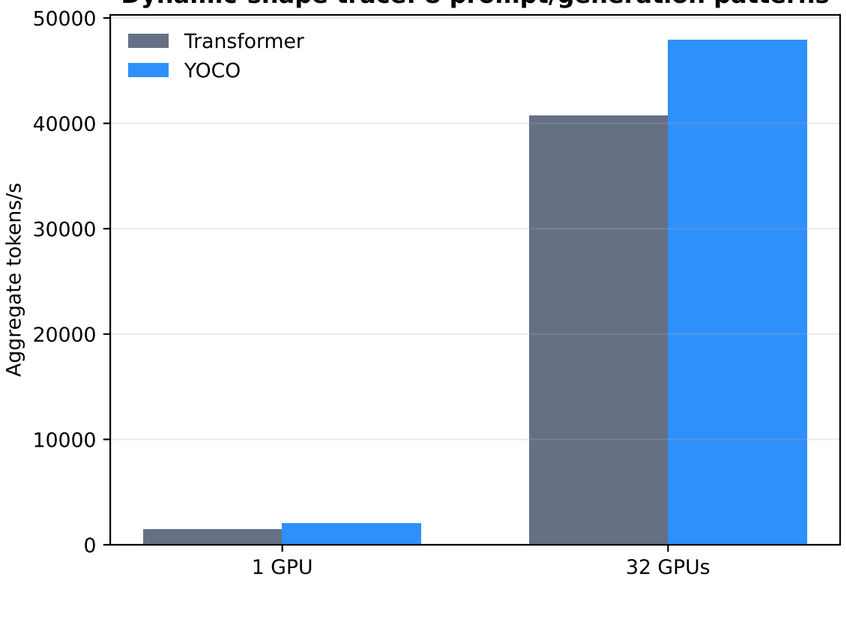

32-GPU dynamic-shape end-to-end speedup

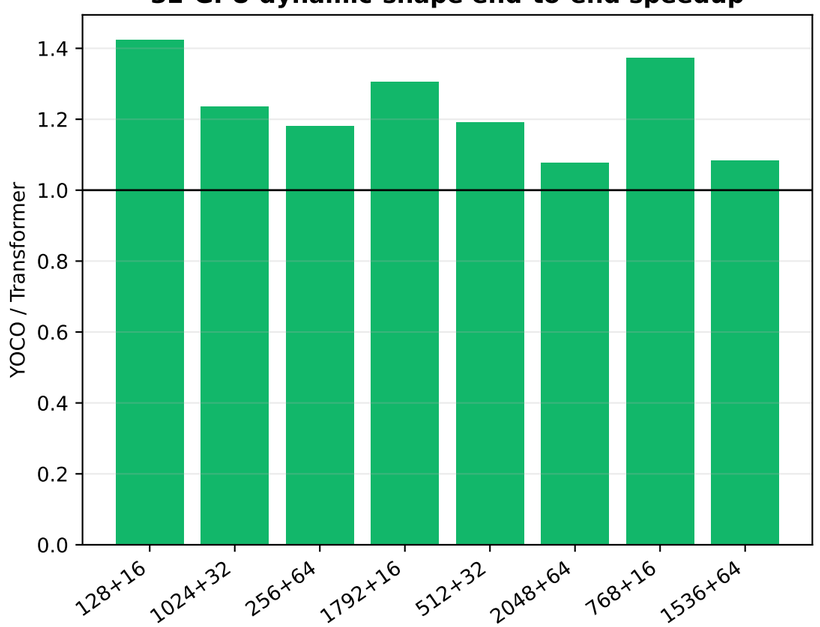

Fixed 2K+64 memory at 32 GPUs

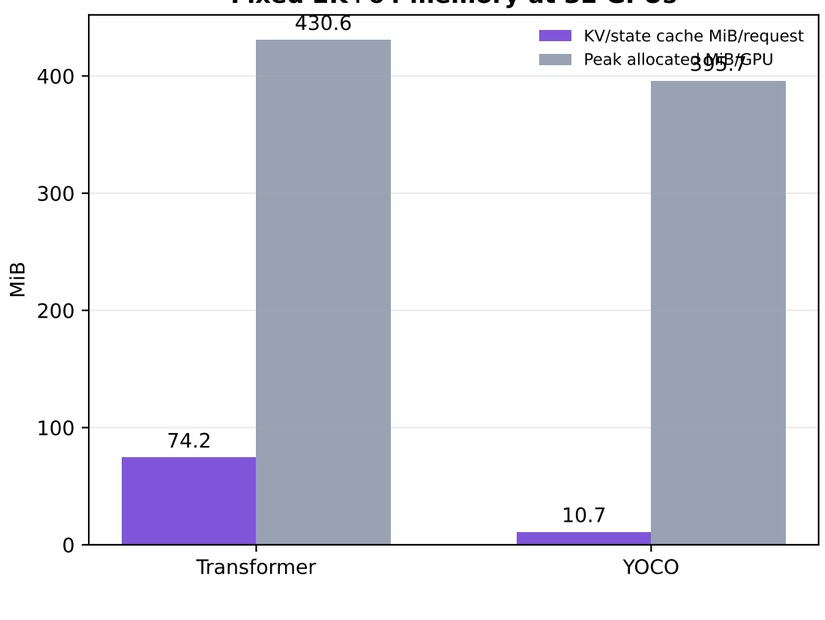

FP16, batch/GPU=1; latency uses slowest rank. 32 GPUs = 4 nodes × 8 V100, replicated serving.